

9.2 — Value categories (lvalues and rvalues)

ALEX JANUARY 18, 2022

Before we talk about our first compound type (lvalue references), we're going to take a little detour and talk about what an **lvalue** is.

In lesson 1.10 — Introduction to expressions, we defined an expression as, "a combination of literals, variables, operators, and function calls that can be executed to produce a singular value". For example:

```
1 #include <iostream>
2
3 int main()
4 {
5     std::cout << 2 + 3; // The expression 2 + 3 produces the value 5
6     return 0;
7 }
```

In the above program, the expression `2 + 3` is evaluated to produce the value 5, which is then printed to the console.

In lesson 5.4 — Increment/decrement operators, and side effects, we also noted that expressions can produce side effects that outlive the expression:

```
1 #include <iostream>
2
3 int main()
4 {
5     int x { 5 };
6     ++x; // This expression statement has the side-effect of incrementing x
7     std::cout << x; // prints 6
8     return 0;
9 }
```

In the above program, the expression `++x` increments the value of `x`, and that value remains changed even after the expression has finished evaluating. Besides producing values and side effects, expressions can do one more thing: they can evaluate to objects or functions. We'll explore this point further in just a moment.

The properties of an expression

To help determine how expressions should evaluate and where they can be used, all expressions in C++ have two properties: a type and a value category.

The type of an expression

The type of an expression is equivalent to the type of the value, object, or function that results from the evaluated expression. For example:

```
1 #include <iostream>
2
3 int main()
4 {
5     auto v1 { 12 / 4 }; // int / int => int
6     auto v2 { 12.0 / 4 }; // double / int => double
7     return 0;
8 }
```

For `v1`, the compiler will determine (at compile time) that a division with two `int` operands will produce an `int` result, so `int` is the type of this expression. Via type inference, `int` will then be used as the type of `v1`.

For `v2`, the compiler will determine (at compile time) that a division with a `double` operand and an `int` operand will produce a `double` result. Remember that arithmetic operators must have operands of matching types, so in this case, the `int` operand gets promoted to a `double`, and a floating point division is performed. So `double` is the type of this expression.

The compiler can use the type of an expression to determine whether an expression is valid in a given context. For example:

```
1 #include <iostream>
2
3 void print(int x)
4 {
5     std::cout << x;
6 }
7
8 int main()
9 {
10    print("foo"); // error: print() was expecting an int argument, we tried to pass in a string literal
11    return 0;
12 }
```

In the above program, the `print(int)` function is expecting an `int` parameter. However, the type of the expression we're passing in (the string literal `"foo"`) does not match, and no conversion can be found. So a compile error results.

Note that the type of an expression must be determinable at compile time (otherwise type checking and type deduction wouldn't work) — however, the value of an expression may be determined at either compile time (if the expression is constexpr) or runtime (if the expression is not constexpr).

The value category of an expression

Now consider the following program:

```
1 int main()
2 {
3     int x{};
4
5     x = 5; // valid: we can assign 5 to x
6     5 = x; // error: can not assign value of x to literal value 5
7     return 0;
8 }
```

One of these assignment statements is valid (assigning value 5 to variable `x`) and one is not (what would it mean to assign the value of `x` to the literal value 5?). So how does the compiler know which expressions can legally appear on either side of an assignment statement?

The answer lies in the second property of expressions: the **value category**. The **value category** of an expression indicates whether an expression resolves to a value, a function, or an object of some kind.

Prior to C++11, there were only two possible value categories: **lvalue** and **rvalue**.

In C++11, three additional value categories (**glvalue**, **prvalue**, and **xvalue**) were added to support a new feature called **move semantics**.

Author's note

In this lesson, we'll stick to the pre-C++11 view of value categories, as this makes for a gentler introduction to value categories (and is all that we need for the moment). We'll cover move semantics (and the additional three value categories) in a future chapter.

Lvalue and rvalue expressions

An **lvalue** (pronounced "ell-value", short for "left value" or "locator value", and sometimes written as "l-value") is an expression that evaluates to a function or object that has an identity. An object or function has an **identity** if it has an identifier (such as a variable or named function) or an identifiable memory address (one that can be retrieved using **operator&**, which we cover in lesson 9.6 — Introduction to pointers). Identifiable objects persist beyond the scope of the expression.

```
1 #include <iostream>
2
3 int main()
4 {
5     int x{};
6
7     std::cout << x << '\n'; // x is an lvalue expression
8     return 0;
9 }
```

In the above program, the expression `x` is an lvalue expression as it evaluates to variable `x` (which has an identifier).

Since the introduction of constants into the language, lvalues come in two subtypes: a **modifiable lvalue** is an lvalue whose value can be modified. A **non-modifiable lvalue** is an lvalue whose value can't be modified (because the lvalue is `const` or `constexpr`).

```
1 #include <iostream>
2
3 int main()
4 {
5     int x{};
6     const double d{};
7
8     std::cout << x << '\n'; // x is a modifiable lvalue expression
9     std::cout << d << '\n'; // d is a non-modifiable lvalue expression
10    return 0;
11 }
```

An **rvalue** (pronounced "arr-value", short for "right value", and sometimes written as **r-value**) is an expression that is not an l-value. Commonly seen rvalues include literals (except string literals, which are lvalues) and the return value of functions or operators. Rvalues only exist within the scope of the expression in which they are used.

```
1 #include <iostream>
2
3 int return5()
4 {
5     return 5;
6 }
7
8 int main()
9 {
10    int x { 5 }; // 5 is an rvalue expression
11    const double d { 1.2 }; // 1.2 is an rvalue expression
12
13    std::cout << x << '\n'; // x is a modifiable lvalue expression
14    std::cout << d << '\n'; // d is a non-modifiable lvalue expression
15    std::cout << return5(); // return5() is an rvalue expression (since the result is returned by value)
16    std::cout << x + 1 << '\n'; // x + 1 is a rvalue
17    std::cout << static_cast<int>(d) << '\n'; // the result of static casting d to an int is an rvalue
18    return 0;
19 }
```

You may be wondering why `return5()` and `x + 1` are rvalues: the answer is because these expressions produce values that must be used immediately (within the scope of the expression) or they are discarded.

Now we can answer the question about why `x = 5` is valid but `5 = x` is not: an assignment operation requires the left operand of the assignment to be a modifiable lvalue expression, and the right operand to be an rvalue expression. The latter assignment (`5 = x`) fails because the expression `5` isn't an lvalue.

```
1 int main()
2 {
3     int x{};
4
5     // Assignment requires the left operand to be a modifiable lvalue expression and the right operand to be an rvalue expression
6     x = 5; // valid: x is a modifiable lvalue expression and 5 is an rvalue expression
7     5 = x; // error: 5 is an rvalue expression and x is a modifiable lvalue expression
8     return 0;
9 }
```

Related content

A full list of lvalue and rvalue expressions can be found [here](#). In C++11, rvalues are broken into two subtypes: prvalues and xvalues, so the rvalues we're talking about here are the sum of both of those categories.

L-value to r-value conversion

We said above that the assignment operator expects the right operand to be an rvalue expression, so why does code like this work?

```
1 int main()
2 {
3     int x { 1 };
4     int y { 2 };
5
6     x = y; // y is a modifiable lvalue, not an rvalue, but this is legal
7     return 0;
8 }
```

The answer is because lvalues will implicitly convert to rvalues, so an lvalue can be used wherever an rvalue is required.

Now consider this snippet:

```
1 int main()
2 {
3     int x { 2 };
4
5     x = x + 1;
6     return 0;
7 }
```

In this statement, the variable `x` is being used in two different contexts. On the left side of the assignment operator, `x` is an lvalue expression that evaluates to variable `x`. On the right side of the assignment operator, `x + 1` is an rvalue expression that evaluates to the value 3.

Now that we've covered lvalues, we can get to our first compound type: the **lvalue reference**.

Key insight

As a rule of thumb to identify lvalue and rvalue expressions:

lvalues expressions are those that evaluate to variables or other identifiable objects that persist beyond the end of the expression.
rvalues expressions are those that evaluate to literals or the returned value of functions and operators that are discarded at the end of the expression.

Next lesson
9.3 Lvalue references

Back to table of contents

Previous lesson
9.1 Introduction to compound data types

B U URL INLINE CODE C++ CODE BLOCK HELP

Leave a comment...



Name*

Notify me about replies:



POST COMMENT



Email*



Avatars from https://gravatar.com/ are connected to your provided email address.

4 COMMENTS

Newest



cheems

February 3, 2022 8:25 am

"r values are the values discarded at the end of expression" i couldn't understand it.
p.s: sorry for asking a noob question

1

Reply



Alex

Author

Reply to cheems February 3, 2022 4:28 pm

rvalues are expressions that evaluate to temporary values. They are typically only needed for a single use, so they are evaluated and then we don't need them any more, so they can be thrown away.

1

Reply



Gabriel

January 23, 2022 1:31 pm

So lvalue is like a container and rvalue is like the stuff you put into the container ... And once you done rvalue is gone since it "merged" to lvalue

Last edited 1 month ago by Gabriel

0

Reply



Alex

Author

Reply to Gabriel January 24, 2022 11:54 am

That analogy doesn't resonate for me.

lvalues tend to be persistent objects that outlive the expression, and rvalues tend to be temporary values (literals, temporary objects, etc...) that exist only for the scope of the expression.

1

Reply